Machine Learning Foundations: From Geometry to Classification

Akshay Redekar

January 25, 2026

Contents

1 Introduction: Machine Learning Context **3**
1.1 Module Overview . . . 3
1.2 Class Structure and Expectations . . . 3
1.3 Why Study Machine Learning? . . . 3

2 Classification: The Fruit Sorting Problem **4**
2.1 Problem Statement . . . 4
2.2 Feature Selection . . . 4
2.3 Essential Terminology . . . 4
2.4 The Challenge of Manual Rules . . . 5

3 Geometric Visualization of Classification **6**
3.1 The Cartesian Plane Representation . . . 6
3.2 Finding a Separation Boundary . . . 6
3.3 Core Machine Learning Definitions . . . 7
3.4 Extension: Multiple Classes and Complex Boundaries . . . 7
3.5 Practical Example: IPL Win Prediction . . . 8

4 Mathematical Foundations: The Equation of a Line **9**
4.1 What is a Line Mathematically? . . . 9
4.2 Understanding Slope . . . 9
4.3 Parallel Lines . . . 10
4.4 Perpendicular Lines . . . 11

5 General Form and Higher Dimensions **12**
5.1 General Form of a Line . . . 12
5.2 Extension to Three Dimensions: Planes . . . 12
5.3 Hyperplanes in n-Dimensions . . . 13

6 Vector Notation and Linear Algebra **14**
6.1 Vectors: A More Compact Notation . . . 14
6.2 The Dot Product (Inner Product) . . . 14
6.3 Vector Form of the Hyperplane Equation . . . 15

7 Vector Magnitude and Distance **16**
7.1 Norm (Magnitude) of a Vector 16
7.2 Angle Between Vectors 16
7.3 Summary: Progression of Equations 17

8 Half Spaces and Classification **18**
8.1 Dividing Space with a Hyperplane 18

9 The "Best" Decision Boundary **20**
9.1 What Makes a Boundary "Best"? 20
9.2 The Margin Principle 20
9.3 Distance Computation Preview 21

10 Forms of Line Equations: Complete Reference **22**
10.1 Five Standard Forms 22

11 Comprehensive Worked Examples **23**
11.1 Example 1: Complete Classification Pipeline 23
11.2 Example 2: Understanding Misclassification 23
11.3 Example 3: IPL Match Prediction 24

12 Practice Problems and Exercises **25**

13 Summary and Next Steps **27**
13.1 Key Concepts Covered 27
13.2 Practical Skills to Develop 28
13.3 Connection to Real ML Algorithms 28

1 Introduction: Machine Learning Context

1.1 Module Overview

This module represents the last foundations component before building your first machine learning model. The primary goal is to understand the mathematical principles that enable machines to learn from data. The module spans approximately 10 lectures over one month and covers four core mathematical topics essential for machine learning.

Core Topics:

- Coordinate Geometry - Understanding how to represent data spatially
- Linear Algebra - Vector operations and transformations
- Calculus - Optimization through derivatives and gradients
- Optimization - Finding the best model parameters

Learning Flow: The pedagogy follows a systematic progression from abstract concepts to concrete implementations. Each topic moves through: Concepts → Visualizations → Mathematics → Code. This approach ensures both theoretical understanding and practical implementation skills.

1.2 Class Structure and Expectations

Class 1 Approach: The first class takes a deliberate, measured pace. It introduces machine learning through practical examples and emphasizes geometric intuition before diving into formal mathematics. Key focus areas include visualization techniques and fundamental equations, particularly the equation of a line which serves as the foundation for linear decision boundaries.

Subsequent Classes: Starting from Class 2, the material becomes more rigorous with heavier emphasis on mathematical formalism and geometric interpretations. This progression allows students to build intuition first before tackling more complex mathematical frameworks.

1.3 Why Study Machine Learning?

Machine learning enables computers to make decisions and predictions without being explicitly programmed with every possible rule. Instead of writing thousands of conditional statements, we can train models to discover patterns in data automatically. This capability has applications ranging from industrial automation to predictive analytics.

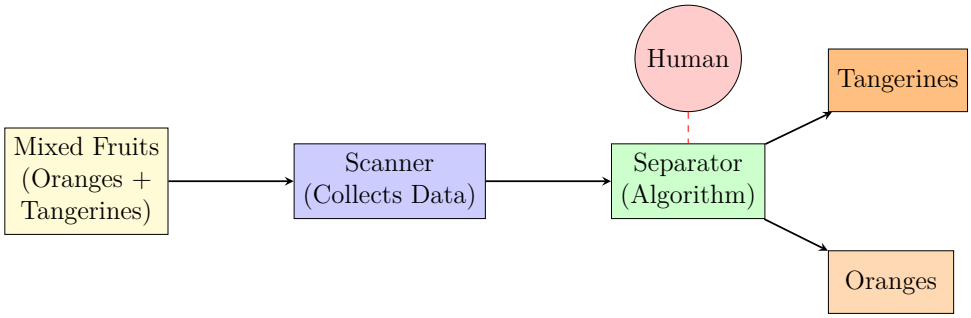
Real-World Applications Include:

- Quality control and sorting in manufacturing
- Predictive modeling in sports and entertainment
- Medical diagnosis and healthcare analytics
- Financial forecasting and risk assessment
- Natural language processing and computer vision

2 Classification: The Fruit Sorting Problem

2.1 Problem Statement

Consider an assembly line at an orange juice factory that processes both oranges and tangerines. Currently, a human worker examines each fruit passing through a scanner and manually operates a separator to direct fruits into the appropriate bins. Our goal is to replace this manual process with an automated algorithm that can classify each fruit correctly.



Goal: Design an algorithm to replace the human worker by automatically classifying fruits based on measurable attributes.

2.2 Feature Selection

The first critical step in any machine learning problem is identifying which attributes (features) are useful for making predictions. Not all measurable properties are equally informative for classification.

Question: Which attributes might help distinguish oranges from tangerines?

Analysis of Potential Features:

- **Diameter:** (Valid) Useful - Can be collected from the scanner. Oranges and tangerines typically have different size distributions.
- **Weight:** (Valid) Useful - Can be collected from the scanner. Weight correlates with size and density differences between the two fruits.
- **Color:** (Invalid) Not useful - Both fruits are orange-colored, making color a poor discriminator for this specific classification task.

This feature selection process demonstrates an important principle: effective machine learning requires domain knowledge to identify which measurements contain discriminative information.

2.3 Essential Terminology

Before proceeding further, we must establish precise definitions for the key concepts used throughout machine learning.

Features
Features (also called independent variables, attributes, or predictors) are the input variables used to make predictions. In our fruit sorting example, the features are Diameter (D) and Weight (W). Features are denoted as $\mathbf{x} = [x_1, x_2, \ldots, x_d]^T$ where d is the number of features.

Target
The target (also called the dependent variable or label) is the output we want to predict. In our example, the target is the fruit type (F), which can be either Orange (O) or Tangerine (T). The target is typically denoted as y.

Data Point
A data point is a single observation consisting of feature values and the corresponding target value. For example, (10, 80, O) represents a fruit with diameter 10, weight 80, which is an orange.

Dataset
A dataset is a collection of many data points. We typically denote a dataset as $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ where n is the number of examples.

Example Dataset:

Diameter (D)	Weight (W)	Fruit (F)
10	100	T
8	70	O
10	80	O
7	100	T
9	85	O
11	95	T

2.4 The Challenge of Manual Rules

One might initially attempt to solve this classification problem by writing explicit rules. For instance: "If weight $>$ 90, then Tangerine; otherwise Orange" or "If diameter $<$ 8, then Orange; otherwise Tangerine."

Problems with Manual Rules:

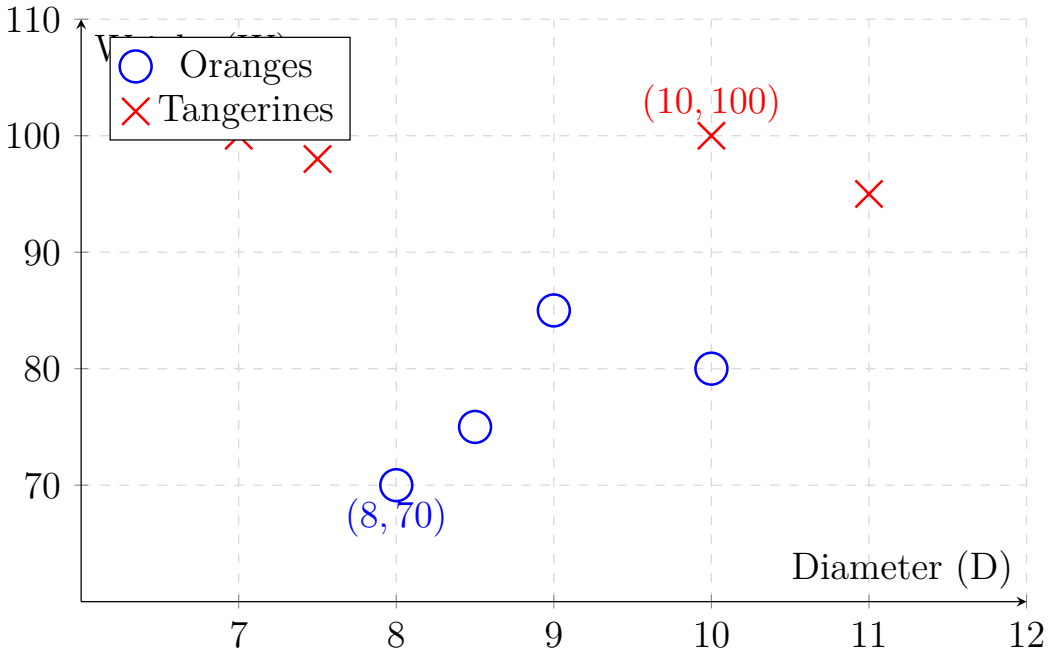
- **Threshold Calibration:** Determining the exact threshold values (like 90 for weight) is difficult and requires extensive trial and error.
- **Multiple Features:** When combining multiple features, the rule complexity grows exponentially. How do we balance diameter and weight together?
- **Rigidity:** Manual rules don't adapt when new data arrives. If fruit sizes change seasonally, rules must be manually updated.
- **Scalability:** With more features (texture, acidity, sugar content), manual rule creation becomes impractical.

This motivates the need for a more systematic, mathematical approach to classification.

3 Geometric Visualization of Classification

3.1 The Cartesian Plane Representation

Rather than thinking in terms of discrete rules, we can visualize the classification problem geometrically. Each data point can be represented as a coordinate in a Cartesian plane where the x-axis represents Diameter (D) and the y-axis represents Weight (W).



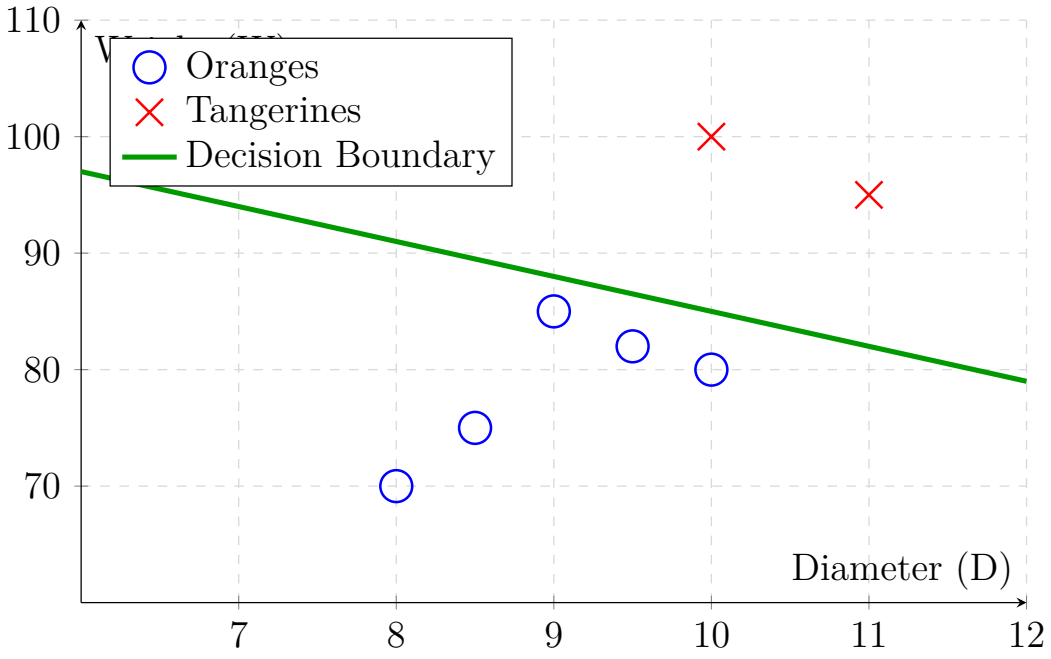
Key Observation: When plotted geometrically, the two classes of fruits form distinct clusters. Oranges tend to occupy one region of the feature space, while tangerines occupy another. This spatial separation is the foundation of classification algorithms.

3.2 Finding a Separation Boundary

Question: Looking at the scatter plot above, can you think of a geometric way to separate these two groups of points?

Answer: A straight line!

We can draw a line that separates the orange points from the tangerine points. Any new fruit that appears can then be classified based on which side of the line it falls on.



3.3 Core Machine Learning Definitions

Decision Boundary
A decision boundary is a geometric surface (line in 2D, plane in 3D, hyperplane in higher dimensions) that separates different classes in the feature space. Points on one side of the boundary are classified as one class, while points on the other side belong to the other class.

Misclassification
A misclassified point is a data point that falls on the wrong side of the decision boundary. For example, if an orange appears in the tangerine region, it represents a classification error.

Machine Learning
Machine learning is the ability of a computer system to automatically find the optimal decision boundary by ”learning” from the given dataset, rather than having the boundary defined manually through hard-coded rules.

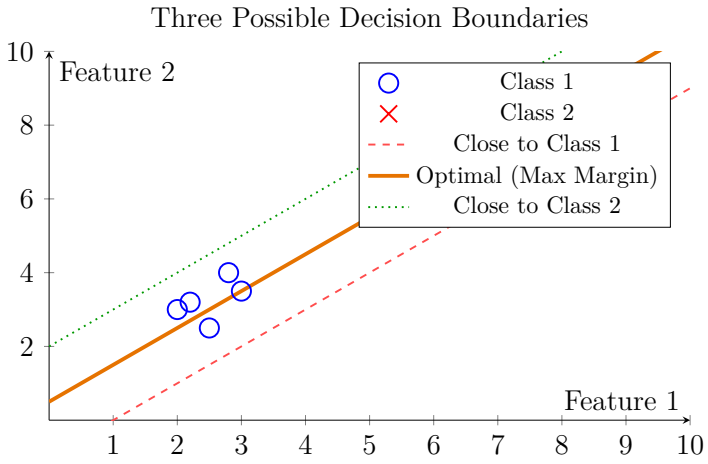
3.4 Extension: Multiple Classes and Complex Boundaries

While our fruit sorting example involves only two classes (binary classification), real-world problems often involve multiple categories. For instance, classifying fruits into apples, oranges, tangerines, and lemons would require multiple decision boundaries.

Important Considerations:

- The number of classes or labels can exceed two, requiring more sophisticated boundary configurations.
- Decision boundaries need not be straight lines - they can be curves, circles, or complex shapes depending on how the data is distributed.

- The choice of boundary shape impacts both model accuracy and computational complexity.



Critical Question: Which line is the "best" decision boundary?

Answer: The orange line represents the optimal boundary because it maintains the maximum distance from both classes. This concept of maximizing the margin between classes is fundamental to support vector machines and robust classification.

3.5 Practical Example: IPL Win Prediction

To further illustrate classification, consider predicting cricket match outcomes for the Indian Premier League (IPL). A marketing company needs to decide when to stop an advertising campaign based on match progress.

Dataset Structure:

- **Features:** Runs scored, Overs completed
- **Target:** Outcome (Win or Loss)

When visualized as a scatter plot with Runs on one axis and Overs on the other, winning matches cluster in one region while losing matches cluster elsewhere. A decision boundary can separate these regions, allowing the system to predict match outcomes in real-time and trigger marketing decisions accordingly.

4 Mathematical Foundations: The Equation of a Line

4.1 What is a Line Mathematically?

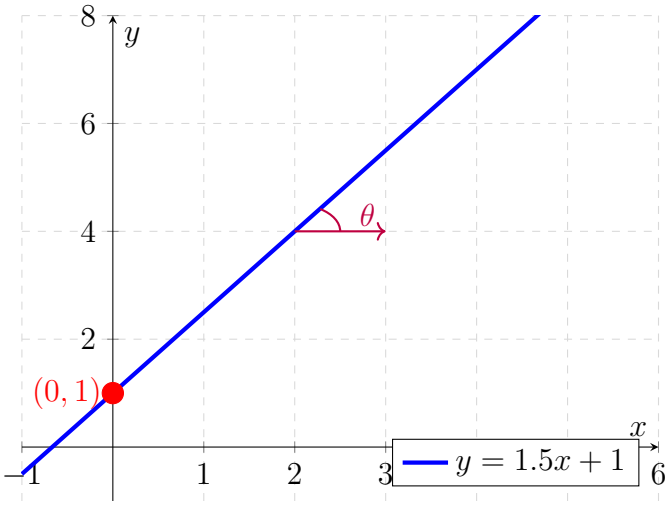
In coordinate geometry, a line represents all points that satisfy a specific linear relationship between coordinates. The most intuitive form is the slope-intercept equation.

Slope-Intercept Form
The slope-intercept form of a line is given by:

$$y = mx + c \tag{1}$$

where:

- m is the slope (gradient) - measures the steepness of the line
- c is the y-intercept - the point where the line crosses the y-axis



4.2 Understanding Slope

The slope quantifies how much y changes for a unit change in x. It can be computed using any two points on the line.

Property 4.1 (Slope Formula)**.** For any two points (x_1, y_1) and (x_2, y_2) on a line:

$$m = \frac{\Delta y}{\Delta x} = \frac{y_2 - y_1}{x_2 - x_1} \tag{2}$$

Relationship to Angle: The slope is also related to the angle θ that the line makes with the positive x-axis:

$$m = \tan(\theta) \quad \text{or equivalently} \quad \theta = \tan^{-1}(m) \tag{3}$$

Example 4.1 (Identifying Linear Relationships)**. Problem 1:** Given points $(1, 2), (2, 4), (3, 6), (4, 8)$, what is the equation of the line?

Solution: Computing the slope using first two points:

$$m = \frac{4 - 2}{2 - 1} = \frac{2}{1} = 2 \tag{4}$$

Finding the y-intercept using point $(1, 2)$:

$$2 = 2(1) + c \implies c = 0 \tag{5}$$

Answer: $y = 2x$

Problem 2: Given points with a flatter slope that pass through $(0, 2)$ and $(1, 3)$, find the equation.

Solution:

$$m = \frac{3-2}{1-0} = 1 \tag{6}$$

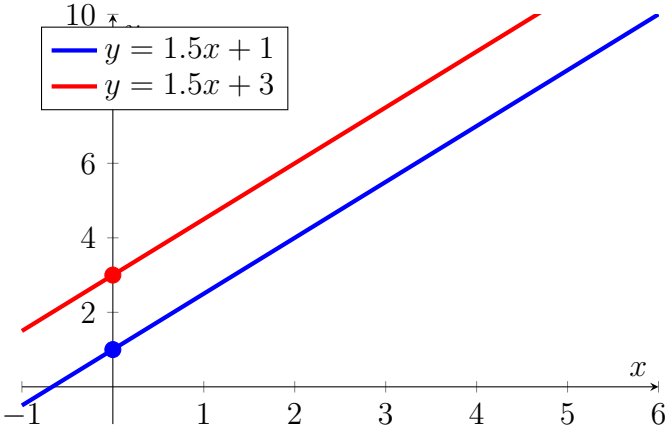
The y-intercept is clearly $c = 2$ (given that the line passes through $(0, 2)$).

Answer: $y = x + 2$

4.3 Parallel Lines

Property 4.2 (Parallel Lines)**.** Two lines are parallel if and only if they have the same slope but different y-intercepts:

$$l_1 \parallel l_2 \iff m_1 = m_2 \text{ and } c_1 \neq c_2 \tag{7}$$



Note: Parallel lines never intersect because they maintain the same angle with the x-axis throughout.

Example 4.2 (Parallel Lines Problem)**. Question:** Find the value of m if the line $y = mx + 3$ is parallel to $0.5y - x = 0$.

Solution: First, convert the second equation to slope-intercept form:

$$0.5y - x = 0 \tag{8}$$

$$0.5y = x \tag{9}$$

$$y = 2x \tag{10}$$

The slope of the second line is $m_2 = 2$. For the lines to be parallel, we need:

$$m_1 = m_2 = 2 \tag{11}$$

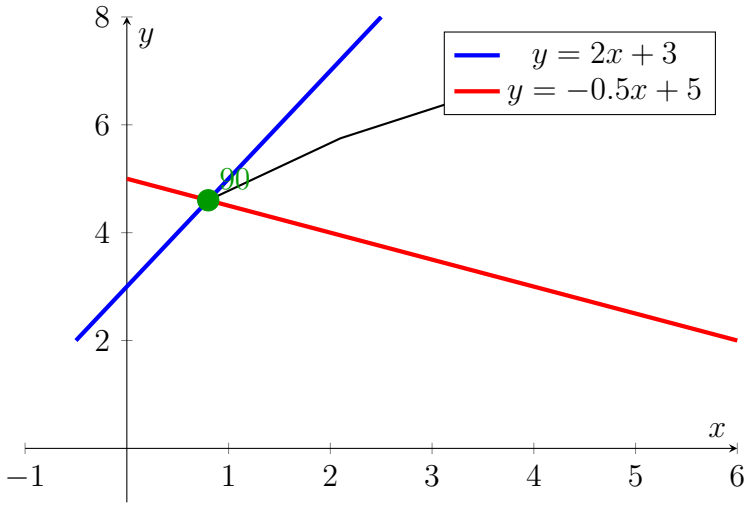
Answer: $m = 2$

4.4 Perpendicular Lines

Property 4.3 (Perpendicular Lines)**.** Two lines are perpendicular if and only if the product of their slopes equals -1:

$$l_1 \perp l_2 \iff m_1 \cdot m_2 = -1 \iff m_1 = -\frac{1}{m_2} \tag{12}$$

Significance for Machine Learning: The distance from a point to a line is measured along the perpendicular direction. Understanding perpendicular lines is therefore crucial for computing classification margins and distances.



Perpendicular Slope Relationship. Consider two lines intersecting at point P. Let θ_1 be the angle that line 1 makes with the x-axis, and θ_2 be the angle for line 2.

For perpendicular lines, the angle between them is 90. This means:

$$\theta_2 = \theta_1 + 90 \tag{13}$$

Using the tangent addition formula:

$$\tan(\theta_1 + 90) = \frac{\tan(\theta_1) + \tan(90)}{1 - \tan(\theta_1)\tan(90)} \tag{14}$$

Since $\tan(90) = \infty$, this simplifies to show that:

$$m_2 = \tan(\theta_2) = -\frac{1}{\tan(\theta_1)} = -\frac{1}{m_1} \tag{15}$$

Therefore: $m_1 \cdot m_2 = -1$ □

5 General Form and Higher Dimensions

5.1 General Form of a Line

While the slope-intercept form $y = mx + c$ is intuitive, the general form provides more flexibility and extends naturally to higher dimensions.

General Form of a Line in 2D
Any line in a 2D plane can be written as:

$$ax + by + c = 0 \tag{16}$$

where a, b, and c are constants, and at least one of a or b is non-zero.

Converting Between Forms:
Starting from slope-intercept form $y = mx + c$:

$$y = mx + c \tag{17}$$
$$mx - y + c = 0 \tag{18}$$
$$mx + (-1)y + c = 0 \tag{19}$$

Conversely, starting from general form $ax + by + c = 0$:

$$ax + by + c = 0 \tag{20}$$
$$by = -ax - c \tag{21}$$
$$y = -\frac{a}{b}x - \frac{c}{b} \tag{22}$$

Therefore:

$$\text{slope } m = -\frac{a}{b}, \quad \text{intercept } c_0 = -\frac{c}{b} \tag{23}$$

Example 5.1 (Converting to Slope-Intercept Form)**. Question:** What is the slope and intercept of $3x + 2y + 5 = 0$?

Solution:

$$3x + 2y + 5 = 0 \tag{24}$$
$$2y = -3x - 5 \tag{25}$$
$$y = -\frac{3}{2}x - \frac{5}{2} \tag{26}$$

Answer: Slope $= -\frac{3}{2}$, Intercept $= -\frac{5}{2}$ (option d)

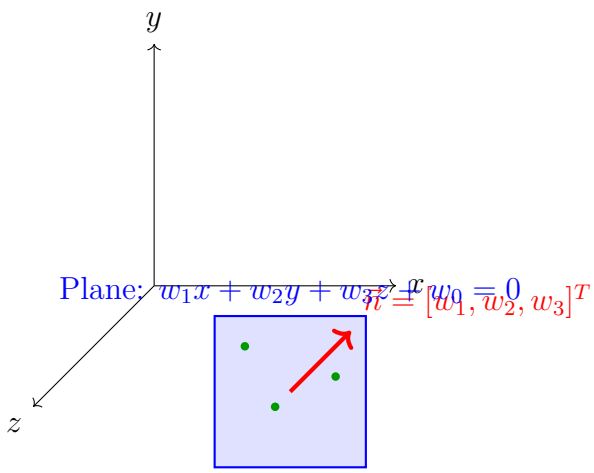
5.2 Extension to Three Dimensions: Planes

When we add a third feature, our decision boundary transitions from a line to a plane.

Equation of a Plane
A plane in 3D space is described by:

$$w_1x + w_2y + w_3z + w_0 = 0 \tag{27}$$

where (x, y, z) are coordinates in 3D space, and w_1, w_2, w_3, w_0 are constants.



5.3 Hyperplanes in n-Dimensions

For machine learning problems with many features, we need to generalize beyond 3D visualization.

Hyperplane
A hyperplane in n-dimensional space is a linear surface defined by:

$$w_1x_1 + w_2x_2 + w_3x_3 + \cdots + w_nx_n + w_0 = 0 \tag{28}$$

where:

- $x_1, x_2, \ldots, x_n$ are the n features
- $w_1, w_2, \ldots, w_n$ are coefficients (also called weights)
- w_0 is the bias term (intercept in higher dimensions)

Dimension of Hyperplane: A hyperplane in n-dimensional space has dimension $n - 1$. For example:

- In 2D ($n = 2$): hyperplane is a line (1D)
- In 3D ($n = 3$): hyperplane is a plane (2D)
- In 100D ($n = 100$): hyperplane is a 99-dimensional surface

Importance in Machine Learning: The hyperplane equation forms the backbone of many fundamental ML algorithms including:

- Linear classifiers
- Support Vector Machines (SVM)
- Logistic regression
- Neural networks (each neuron computes a weighted sum forming a hyperplane)

6 Vector Notation and Linear Algebra

6.1 Vectors: A More Compact Notation

The general hyperplane equation $w_1x_1+w_2x_2+\cdots+w_nx_n+w_0=0$ becomes cumbersome as the number of features increases. Vector notation provides a much more elegant representation.

Vector
A vector is an ordered array of numbers. Conventionally, vectors are written as column vectors:

$$\vec{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}_{n\times 1} \tag{29}$$

The notation $n \times 1$ indicates n rows and 1 column.

Transpose
The transpose of a vector converts it from a column to a row (or vice versa):

$$\vec{v}^T = [v_1, v_2, \ldots, v_n]_{1\times n} \tag{30}$$

The transpose operation is denoted by the superscript T.

Application to Our Problem:
We can represent the feature vector and weight vector as:

$$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, \quad \vec{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix} \tag{31}$$

Both vectors belong to $\mathbb{R}^n$ (the set of all n-dimensional real-valued vectors), while the bias $w_0 \in \mathbb{R}^1$ is a scalar.

6.2 The Dot Product (Inner Product)

Dot Product
The dot product (also called inner product or scalar product) of two vectors $\vec{a}$ and $\vec{b}$ is:

$$\vec{a}\cdot\vec{b} = a_1b_1 + a_2b_2 + a_3b_3 + \cdots + a_nb_n = \sum_{i=1}^{n} a_ib_i \tag{32}$$

Alternative notation: $\langle\vec{a},\vec{b}\rangle$ means the same as $\vec{a}\cdot\vec{b}$

Matrix Multiplication Perspective: The dot product can also be expressed as matrix multiplication:

$$\vec{a}\cdot\vec{b} = \vec{a}^T\vec{b} = [a_1, a_2, \ldots, a_n]\begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix} = a_1b_1 + a_2b_2 + \cdots + a_nb_n \tag{33}$$

Example 6.1 (Computing Dot Products)**. Question:** Find the dot product of $\vec{a} = [1, 2, 3]$ and $\vec{b} = [-1, 2, 5]$.

Solution:

$$\vec{a} \cdot \vec{b} = (1)(-1) + (2)(2) + (3)(5) \tag{34}$$
$$= -1 + 4 + 15 \tag{35}$$
$$= 18 \tag{36}$$

Answer: 18 (option c)

6.3 Vector Form of the Hyperplane Equation

Using vector notation, the lengthy hyperplane equation becomes remarkably concise:

Vector Form of Hyperplane

General Form:

$$w_1x_1 + w_2x_2 + w_3x_3 + \cdots + w_nx_n + w_0 = 0 \tag{37}$$

Vector Form (Equivalent):

$$\vec{w}^T\vec{x} + w_0 = 0 \tag{38}$$

This compact representation is universally used in machine learning implementations.

Expanded View:

$$\vec{w}^T\vec{x} + w_0 = [w_1, w_2, \ldots, w_n]_{1\times n} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}_{n\times 1} + w_0 = 0 \tag{39}$$

The dimensions work out: $(1 \times n) \times (n \times 1) = (1 \times 1)$ which is a scalar, plus the scalar w_0, equals zero.

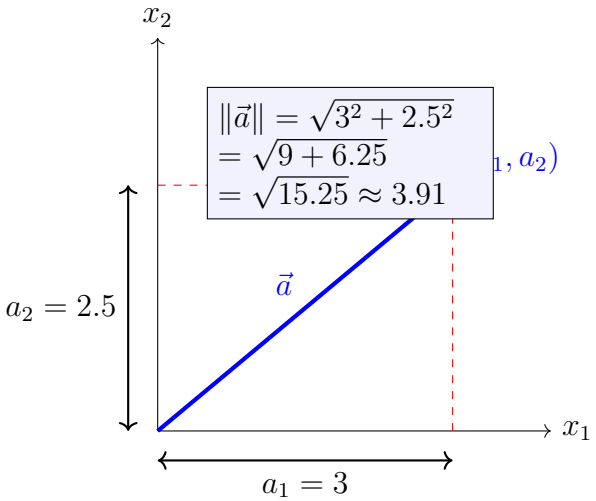
7 Vector Magnitude and Distance

7.1 Norm (Magnitude) of a Vector

Euclidean Norm
The Euclidean norm (also called magnitude or length) of a vector $\vec{a}$ is defined as:

$$\|\vec{a}\| = \sqrt{\vec{a} \cdot \vec{a}} = \sqrt{a_1^2 + a_2^2 + \cdots + a_n^2} \tag{40}$$

This represents the distance from the origin to the point represented by $\vec{a}$.



Example 7.1 (Computing Vector Norms)**. Problem:** Find the norm of $\vec{v} = [3, 4]^T$.
Solution:

$$\|\vec{v}\| = \sqrt{3^2 + 4^2} = \sqrt{9 + 16} = \sqrt{25} = 5 \tag{41}$$

Problem: Find the norm of $\vec{u} = [1, 2, 2]^T$.
Solution:

$$\|\vec{u}\| = \sqrt{1^2 + 2^2 + 2^2} = \sqrt{1 + 4 + 4} = \sqrt{9} = 3 \tag{42}$$

7.2 Angle Between Vectors

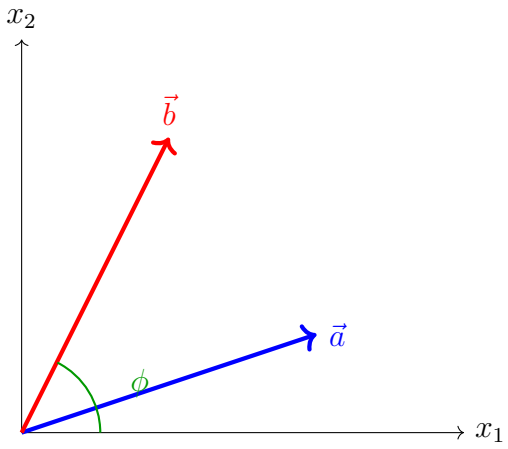
The dot product has a geometric interpretation related to the angle between vectors.

Theorem 7.1 (Dot Product and Angle)**.** For two non-zero vectors $\vec{a}$ and $\vec{b}$, the angle ϕ between them satisfies:

$$\cos(\phi) = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \cdot \|\vec{b}\|} \tag{43}$$

Special Cases:

- If $\vec{a} \cdot \vec{b} = 0$, then $\cos(\phi) = 0$, so $\phi = 90$ (perpendicular vectors)
- If $\vec{a} \cdot \vec{b} > 0$, then $\phi < 90$ (acute angle)
- If $\vec{a} \cdot \vec{b} < 0$, then $\phi > 90$ (obtuse angle)



$\vec{a} = [3, 1]^T, \vec{b} = [1.5, 3]^T$
$\vec{a} \cdot \vec{b} = 3(1.5) + 1(3) = 7.5$
$\|\vec{a}\| = \sqrt{10}, \|\vec{b}\| = \sqrt{11.25}$
$\cos(\phi) = \frac{7.5}{\sqrt{10} \cdot \sqrt{11.25}} \approx 0.707$
$\phi = \cos^{-1}(0.707) \approx 45$

7.3 Summary: Progression of Equations

Dimension	Space	Equation Form	Vector Form
2D	Plane	$y = mx + c$	$ax + by + c = 0$
3D	Space	$w_1x + w_2y + w_3z + w_0 = 0$	$\vec{w}^T\vec{x} + w_0 = 0$
nD	Hyperspace	$\sum_{i=1}^{n} w_ix_i + w_0 = 0$	$\vec{w}^T\vec{x} + w_0 = 0$

Key Insight: The vector form $\vec{w}^T\vec{x} + w_0 = 0$ remains unchanged regardless of dimensionality, making it the preferred notation for machine learning algorithms that must work in arbitrary dimensions.

8 Half Spaces and Classification

8.1 Dividing Space with a Hyperplane

A hyperplane doesn't just sit in space - it actively divides the entire space into two distinct regions called half-spaces.

Half Spaces
For a hyperplane defined by $\vec{w}^T\vec{x} + w_0 = 0$, the space is divided into:

Positive Half Space:

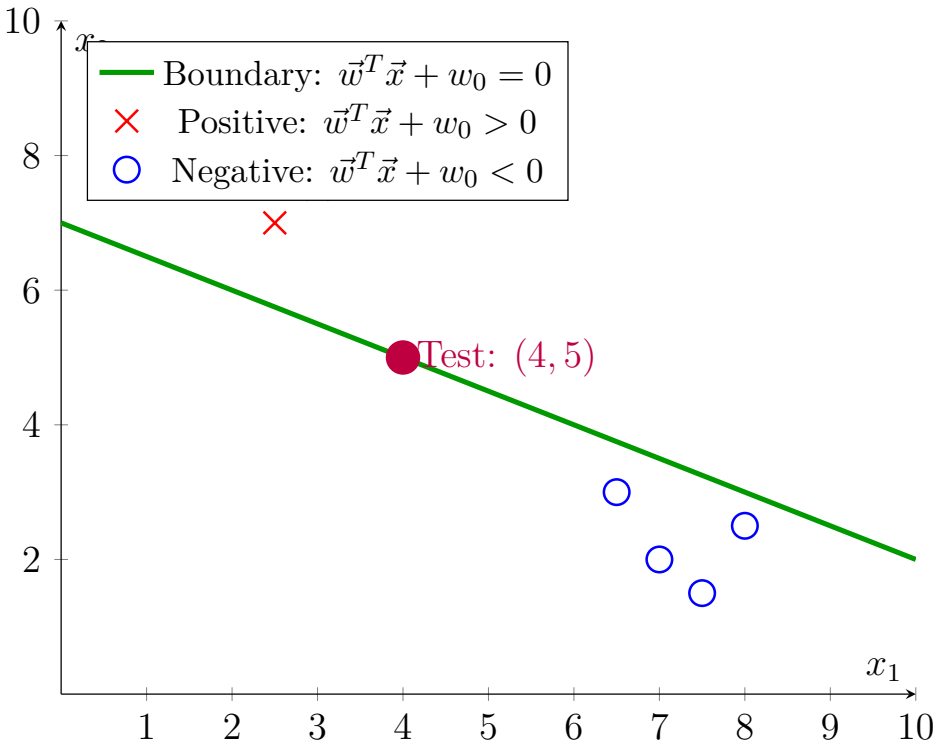
$$\{\vec{x} : \vec{w}^T\vec{x} + w_0 > 0\} \tag{44}$$

Negative Half Space:

$$\{\vec{x} : \vec{w}^T\vec{x} + w_0 < 0\} \tag{45}$$

Classification Using Half Spaces:
For a new point $\vec{x}_{\text{new}}$, we can classify it by computing $\vec{w}^T\vec{x}_{\text{new}} + w_0$:

- If $\vec{w}^T\vec{x}_{\text{new}} + w_0 > 0 \rightarrow$ Positive class (e.g., Tangerine)
- If $\vec{w}^T\vec{x}_{\text{new}} + w_0 < 0 \rightarrow$ Negative class (e.g., Orange)
- If $\vec{w}^T\vec{x}_{\text{new}} + w_0 = 0 \rightarrow$ Exactly on the boundary



Example 8.1 (Half Space Classification)**. Given:** Decision boundary $2x_1 + x_2 - 10 = 0$ (equivalently, $\vec{w} = [2, 1]^T, w_0 = -10$)

Classify: Point $(4, 5)$

Solution:

$$\begin{aligned} \vec{w}^T\vec{x} + w_0 &= 2(4) + 1(5) - 10 && (46)\\ &= 8 + 5 - 10 && (47)\\ &= 3 > 0 && (48) \end{aligned}$$

Since the result is positive, the point $(4, 5)$ belongs to the positive half space.
Classify: Point $(1, 2)$
Solution:

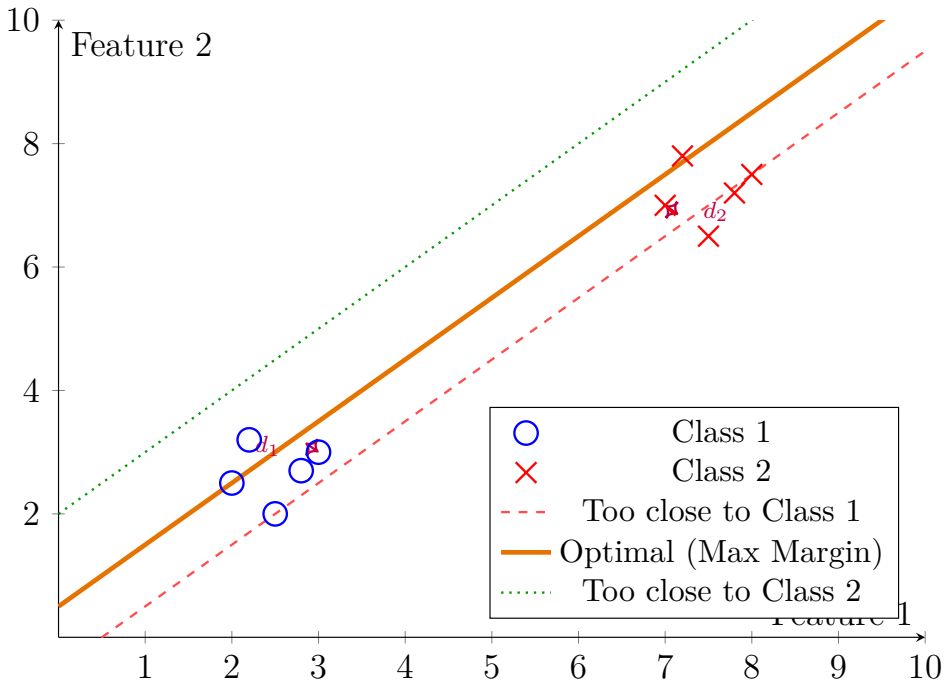
$$\vec{w}^T\vec{x} + w_0 = 2(1) + 1(2) - 10 \tag{49}$$
$$= 2 + 2 - 10 \tag{50}$$
$$= -6 < 0 \tag{51}$$

Since the result is negative, the point $(1, 2)$ belongs to the negative half space.

9 The "Best" Decision Boundary

9.1 What Makes a Boundary "Best"?

We've established that we can draw lines to separate classes, but there are infinitely many lines that could separate two clusters. The crucial question is: which line should we choose?



9.2 The Margin Principle

Optimal Decision Boundary Principle

The **best** decision boundary is the one that maximizes the minimum distance to points from both classes. This is called maximizing the **margin**.
Why? A larger margin provides:

- Greater robustness to small variations in data
- Better generalization to new, unseen examples
- Lower sensitivity to measurement noise
- Improved confidence in predictions near the boundary

Mathematical Formulation:
The optimal boundary $\vec{w}^*$ and w_0^* solve:

$$\max_{\vec{w},w_0} \left\{ \min_{i=1,\ldots,n} \text{distance}(\vec{x}_i, \text{boundary}) \right\} \tag{52}$$

This optimization problem requires us to:

1. Compute distances from points to the boundary
2. Find the configuration that maximizes the smallest such distance

9.3 Distance Computation Preview

To find the optimal boundary, we need a formula for the distance from a point to a hyperplane. This is where our linear algebra tools become essential.

Distance from point $\vec{x}_i$ to hyperplane $\vec{w}^T\vec{x} + w_0 = 0$:

$$d = \frac{|\vec{w}^T\vec{x}_i + w_0|}{\|\vec{w}\|} \tag{53}$$

The derivation of this formula requires understanding:

- Vector projections
- Perpendicular distances
- Geometric properties of dot products

These topics will be covered in depth in subsequent lectures focusing on linear algebra and vector calculus.

10 Forms of Line Equations: Complete Reference

10.1 Five Standard Forms

Different forms of line equations are useful in different contexts. Here's a comprehensive overview:

1. **Slope-Intercept Form:**
$$y = mx + c \tag{54}$$
Use: When slope and intercept are known; most intuitive form.

2. **Point-Slope Form:**
$$y - y_1 = m(x - x_1) \tag{55}$$
Use: When you know one point (x_1, y_1) and the slope m.

3. **Two-Point Form:**
$$y - y_1 = \frac{y_2 - y_1}{x_2 - x_1}(x - x_1) \tag{56}$$
Use: When you know two points (x_1, y_1) and (x_2, y_2) on the line.

4. **Intercept Form:**
$$\frac{x}{a} + \frac{y}{b} = 1 \tag{57}$$
Use: When x-intercept is a and y-intercept is b.

5. **General Form:**
$$ax + by + c = 0 \tag{58}$$
Use: Standard form; generalizes to hyperplanes; used in ML algorithms.

Example 10.1 (Converting Between Forms)**. Given:** Two points $(2, 3)$ and $(4, 7)$
Find: All five forms of the line equation.
Solution:
Step 1: Find slope
$$m = \frac{7 - 3}{4 - 2} = \frac{4}{2} = 2 \tag{59}$$
Step 2: Find y-intercept using point $(2, 3)$
$$3 = 2(2) + c \implies c = -1 \tag{60}$$
Forms:

1. Slope-intercept: $y = 2x - 1$
2. Point-slope: $y - 3 = 2(x - 2)$ or $y - 7 = 2(x - 4)$
3. Two-point: $y - 3 = \frac{7-3}{4-2}(x - 2) = 2(x - 2)$
4. General: $2x - y - 1 = 0$
5. Intercept: $\frac{x}{0.5} + \frac{y}{-1} = 1$ (x-intercept is 0.5, y-intercept is -1)

11 Comprehensive Worked Examples

11.1 Example 1: Complete Classification Pipeline

Example 11.1 (End-to-End Fruit Classification)**. Given Dataset:**

Diameter	Weight	Fruit
7	65	O
8	75	O
9	80	O
10	95	T
11	100	T
12	105	T

Task: Find a linear decision boundary to separate oranges from tangerines.

Solution Approach:

Step 1: Visualize the data

Plot the points on a Diameter-Weight plane. Visual inspection suggests a line with negative slope could separate the classes.

Step 2: Propose a boundary

Let's try the line: $W = -3D + 115$ or in general form: $3D + W - 115 = 0$

In vector notation with $\vec{x} = [D, W]^T$:

$$\vec{w}^T\vec{x} + w_0 = 3D + 1W - 115 = 0 \tag{61}$$

So $\vec{w} = [3, 1]^T$ and $w_0 = -115$.

Step 3: Test all points

For Oranges (should be in negative half space):

$$(7, 65): \quad 3(7) + 65 - 115 = -29 < 0 \quad (Correct) \tag{62}$$
$$(8, 75): \quad 3(8) + 75 - 115 = -16 < 0 \quad (Correct) \tag{63}$$
$$(9, 80): \quad 3(9) + 80 - 115 = -8 < 0 \quad (Correct) \tag{64}$$

For Tangerines (should be in positive half space):

$$(10, 95): \quad 3(10) + 95 - 115 = 10 > 0 \quad (Correct) \tag{65}$$
$$(11, 100): \quad 3(11) + 100 - 115 = 16 > 0 \quad (Correct) \tag{66}$$
$$(12, 105): \quad 3(12) + 105 - 115 = 26 > 0 \quad (Correct) \tag{67}$$

Result: This boundary perfectly separates all training points!

11.2 Example 2: Understanding Misclassification

Example 11.2 (Handling Imperfect Boundaries)**. Modified Dataset:** Add a noisy point: $(9, 98, O)$ - an unusually heavy orange.

Analysis with previous boundary $3D + W - 115 = 0$**:**

$$3(9) + 98 - 115 = 27 + 98 - 115 = 10 > 0 \tag{68}$$

This orange would be misclassified as a tangerine because it falls in the positive half space!

Options:

1. **Adjust the boundary:** Find a new line that correctly classifies this point (may misclassify others)
2. **Accept the error:** Recognize that perfect separation isn't always possible with linear boundaries
3. **Use complex boundaries:** Consider non-linear decision boundaries (covered in advanced ML)

This example illustrates that real-world classification often involves trade-offs between different types of errors.

11.3 Example 3: IPL Match Prediction

Example 11.3 (Cricket Match Outcome Prediction)**. Scenario:** A marketing company runs ads during IPL matches. They want to stop advertising early if a match is already decided to save costs.

Features:

- x_1: Runs scored by batting team
- x_2: Overs completed

Target: Win (W) or Loss (L)
Sample Data:

Runs	Overs	Outcome
45	6	L
78	10	W
52	8	L
95	12	W
110	15	W
60	10	L

Visualization Insight: When plotted, winning matches (W) cluster in the upper-right region (high runs, moderate overs), while losing matches (L) cluster in the lower-left (low runs). A decision boundary can separate these regions.

Business Application: Once the boundary is learned, the system can predict in real-time: "If current position is in the win region with high confidence, continue ads; if in loss region, stop ads to save budget."

12 Practice Problems and Exercises

Problem 12.1 (Line Equations)**.** For each set of information, write the equation of the line in both slope-intercept and general form:

(a) Slope $m = 3$, y-intercept $c = -2$

(b) Passes through $(1, 4)$ and $(3, 10)$

(c) x-intercept $= 5$, y-intercept $= -3$

(d) Parallel to $2x + 3y - 7 = 0$ and passes through $(0, 4)$

Problem 12.2 (Parallel and Perpendicular Lines)**.** (a) Find the equation of a line parallel to $y = 4x - 1$ passing through $(2, 5)$

(b) Find the equation of a line perpendicular to $y = -\frac{1}{3}x + 2$ passing through the origin

(c) Determine whether the lines $2x - 3y + 5 = 0$ and $3x + 2y - 1 = 0$ are parallel, perpendicular, or neither

Problem 12.3 (Vector Operations)**.** Given $\vec{a} = [2, -1, 3]^T$, $\vec{b} = [1, 4, -2]^T$, and $\vec{c} = [-3, 0, 5]^T$:

(a) Compute $\vec{a} \cdot \vec{b}$

(b) Compute $\|\vec{a}\|$ and $\|\vec{b}\|$

(c) Find the angle between $\vec{a}$ and $\vec{b}$

(d) Verify that $(\vec{a} + \vec{b}) \cdot \vec{c} = \vec{a} \cdot \vec{c} + \vec{b} \cdot \vec{c}$

Problem 12.4 (Half Space Classification)**.** Consider the decision boundary $2x_1 - 3x_2 + 6 = 0$.

(a) Classify the point $(3, 4)$

(b) Classify the point $(0, 0)$

(c) Find a point that lies exactly on the boundary

(d) If class A is in the positive half space and class B is in the negative half space, which class does $(5, 2)$ belong to?

Problem 12.5 (Feature Space Visualization)**.** You are building a classifier to predict if a student will pass or fail based on hours studied and sleep hours.

Dataset:

Study Hours	Sleep Hours	Result
2	4	Fail
5	7	Pass
3	5	Fail
7	6	Pass
8	8	Pass
1	3	Fail

(a) Sketch the scatter plot

(b) Propose a reasonable linear decision boundary

(c) Write the equation of your boundary in vector form

(d) Classify a student with 4 study hours and 6 sleep hours

Problem 12.6 (Multi-class Extension)**.** Suppose we want to classify fruits into three categories: Small (S), Medium (M), Large (L) based on diameter and weight.

(a) How many decision boundaries would we need minimum?

(b) Sketch a possible configuration of boundaries

(c) What happens if the boundaries intersect? Is this problematic?

13 Summary and Next Steps

13.1 Key Concepts Covered

Module 1 Summary

1. **Machine Learning Fundamentals:**
   - Machine learning enables automated decision-making from data
   - Classification involves finding decision boundaries between classes
   - Geometric visualization helps build intuition before formal math
2. **Feature Engineering:**
   - Features are the measurable attributes used for prediction
   - Not all features are equally informative
   - Feature selection requires domain knowledge
3. **Linear Boundaries:**
   - Lines in 2D, planes in 3D, hyperplanes in nD
   - Multiple equivalent forms: slope-intercept, general, vector
   - Vector form $\vec{w}^T\vec{x} + w_0 = 0$ generalizes to any dimension
4. **Optimization Goal:**
   - Best boundary maximizes margin between classes
   - Requires computing distances to hyperplanes
   - Machine learning algorithms automate this search
5. **Linear Algebra Tools:**
   - Vectors provide compact notation for high-dimensional data
   - Dot products enable distance and angle calculations
   - Norms measure vector magnitudes

Example 13.1 (Calculating Distance to Boundary)**. Problem:** Calculate the perpendicular distance from point $P(2, 4)$ to the decision boundary line $3x + 4y - 12 = 0$.

Solution Steps:

1. Identify the parameters from the line equation $ax+by+c = 0$: $a = 3, \quad b = 4, \quad c = -12$
2. Identify the point coordinates (x_0, y_0): $x_0 = 2, \quad y_0 = 4$
3. Apply the distance formula: $d = \frac{|ax_0+by_0+c|}{\sqrt{a^2+b^2}}$

4. Substitute the values:

$$\begin{aligned} d &= \frac{|3(2) + 4(4) - 12|}{\sqrt{3^2 + 4^2}} \\ &= \frac{|6 + 16 - 12|}{\sqrt{9 + 16}} \\ &= \frac{|10|}{\sqrt{25}} \\ &= \frac{10}{5} = 2 \end{aligned}$$

Answer: The distance is 2 units.

13.2 Practical Skills to Develop

To reinforce these concepts, practice the following:

1. **Visualization:** Use Desmos or Python (matplotlib) to plot datasets and experiment with different decision boundaries
2. **Vector Calculations:** Become fluent with dot products, norms, and vector arithmetic
3. **Form Conversion:** Practice converting between different line equation forms quickly
4. **Geometric Intuition:** Train yourself to visualize high-dimensional problems using 2D/3D analogies

13.3 Connection to Real ML Algorithms

The concepts introduced here directly map to production machine learning algorithms:

Concept	**ML Algorithm**
Linear decision boundary	Linear SVM, Perceptron, Logistic Regression
Maximizing margin	Support Vector Machines (SVM)
Half-space classification	All linear classifiers
Hyperplane in nD	Neural network layers (each neuron is a hyperplane)
Finding optimal $\vec{w}$	Gradient descent, backpropagation